The Agent Is a Loop

Author: Konrad Kowalczyk
Repository: ai-agents-JAVA-SPRING

Let's start from the beginning

This is a continuation of last year's presentation, which focused on…



OLD REPO → xkondix AI Sandbox

Roadmap

1. The Agent Is a Loop

Writing one by hand, and what frameworks hide when they write it for you

2. Agentic Workflow Patterns

Five shapes, two frameworks, and when the DSL earns its keep

3. MCP as a Process Boundary

Tools, resources and prompts over two transports: HTTP for agents, STDIO for Claude Desktop

4. How Do You Know It Works?

Traces, metrics, logs and dashboards for everything your agents actually do

The Agent Is a Loop

1. **What is an agent?**
2. **Build it** → No framework
3. **Hide it** → LangChain4j
4. **Control it** → Spring AI
5. **Mix it** → One provider, ∞ protocols



What is an agent?



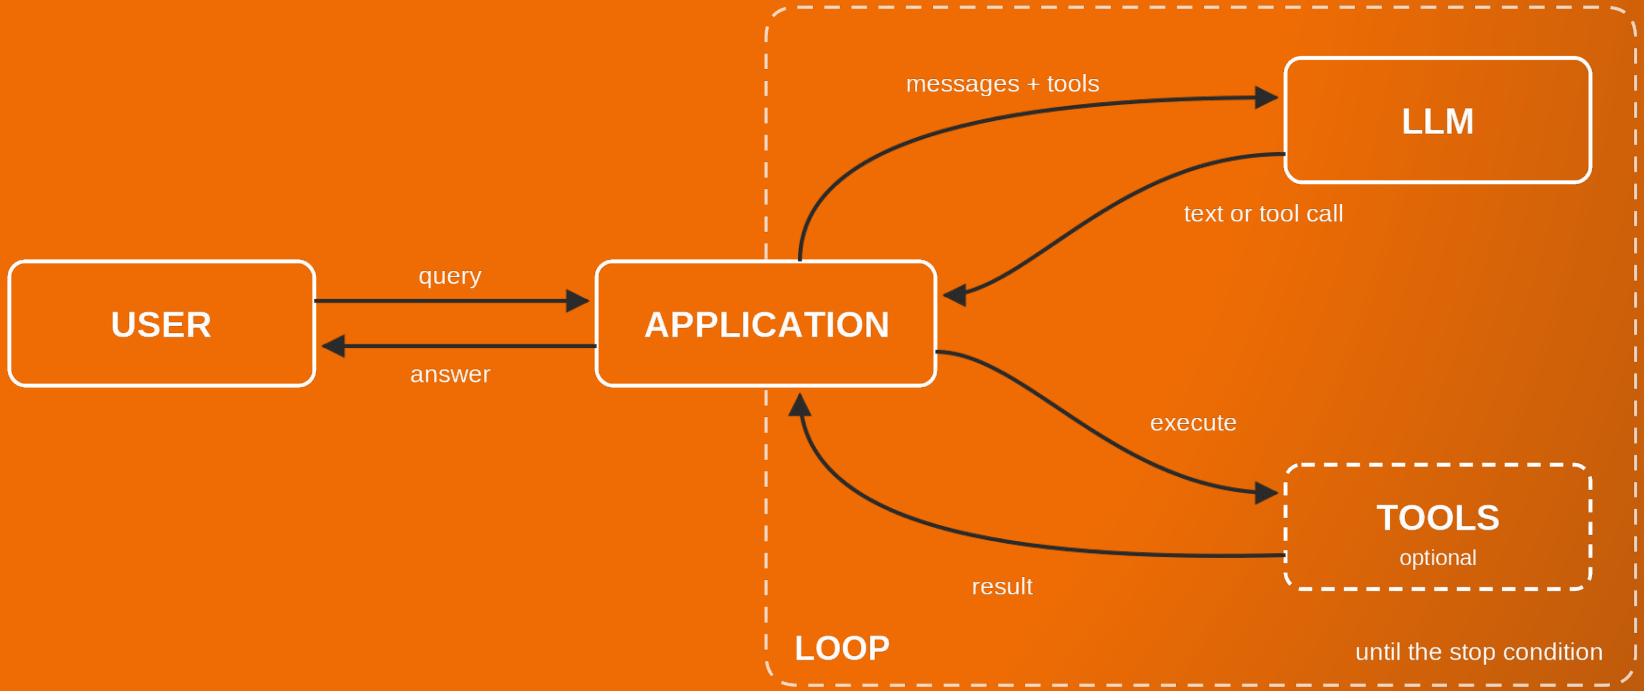
The Agent

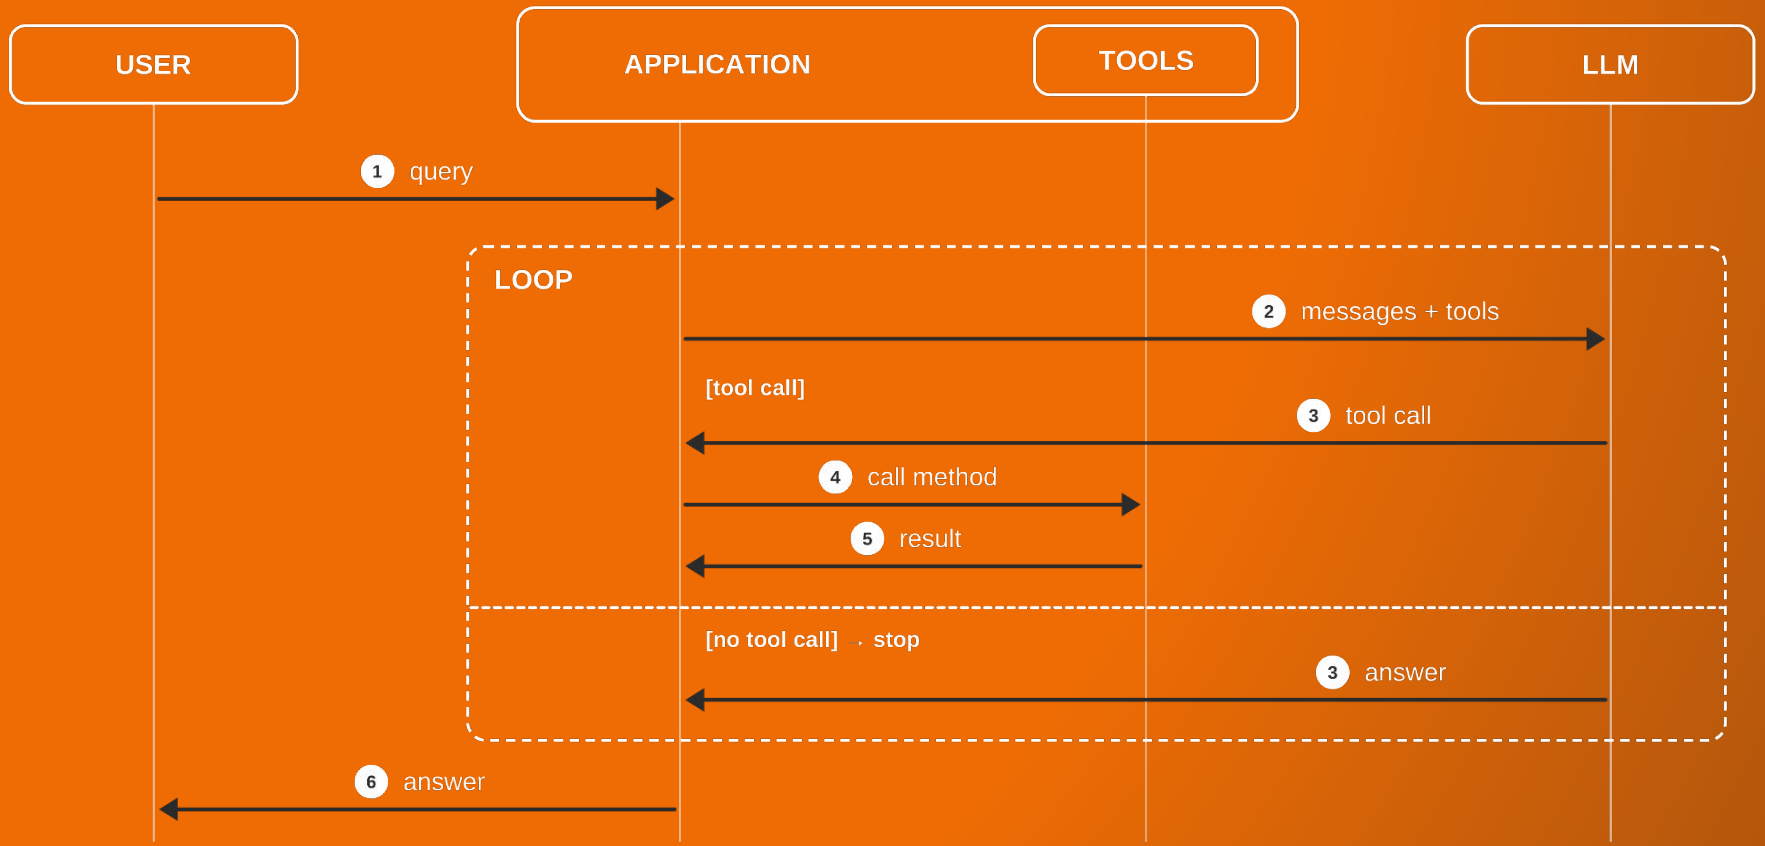


The true Agent

Agents are typically just LLMs using tools based on environmental feedback in a loop.




messages + tools
LLM
text or tool call
query
USER
APPLICATION
answer
execute
TOOLS
optional
result
LOOP
until the stop condition


USER
APPLICATION
TOOLS
LLM
1 query
LOOP
2 messages + tools
[tool call]
3 tool call
4 call method
5 result
[no tool call] → stop
3 answer
6 answer

Raw Agent

The loop exits when the model stops requesting tools.

stop condition = no more tool calls

sources

- openai
- strandsagents

```
// ── Agent loop ──────────────────────────────────────────
for (int i = 1; i <= MAX_ITERATIONS; i++) {
    log.info("[LOOP] Iteration {} — {} messages in history", i, history.size());

    // ── Step 2: Call LLM ────────────────────────────────
    var response = llmClient.chat(history, tools.getDefinitions());
    var message  = response.firstMessage();

    if (message == null) {
        log.error("[LOOP] Empty response from LLM");
        return "Error: empty response from model.";
    }

    history.add(Message.assistant(message.content(), message.toolCalls()));

    if (response.usage() != null) {
        log.info("[LOOP] Tokens: input={} output={} total={}",
                response.usage().promptTokens(),
                response.usage().completionTokens(),
                response.usage().totalTokens());
    }

    // ── Step 3: Tool calls? ─────────────────────────────
    if (!response.hasToolCalls()) {
        log.info("[LOOP] Done after {} iteration(s)", i);
        return message.content() != null ? message.content() : "";
    }

    for (ToolCall toolCall : message.toolCalls()) {
        String result = executeToolWithSpan(toolCall, i);
        history.add(Message.tool(toolCall.id(), result));
    }
}

return "Max iterations (" + MAX_ITERATIONS + ") reached.";
```

The request body structure

Every exchange with the model is JSON. The framework only decides who writes it.

```
{
  "model": "openai/gpt-4o-mini",
  "messages": [ … ],
  "tools": [ … ],
  "temperature": 0.3,
  "stream": false
}
```

MODEL → the model that handles this request

MESSAGES → the whole conversation, resent on every call

TOOLS → name, description and parameters of each method

TEMPERATURE → lower: more deterministic
higher: more random

STREAM → false: one complete response
true: chunks as they are generated

The request body as JSON

MESSAGES []

SYSTEM → the instructions

USER → the question

ASSISTANT → a request for a tool

TOOL → the result your code produced

LOOP

ASSISTANT → the answer (response ends the loop)



```
{
  "model" : "openai/gpt-4o-mini",
  "messages" : [ {
    "role" : "system",
    "content" : "You are a helpful assistant. Use available tools when appropriate."
  }, {
    "role" : "user",
    "content" : "How is the weather in Katowice?"
  }, {
    "role" : "assistant",
    "tool_calls" : [ {
      "id" : "call_EIrKGD9Ec5BB3Ljw94tPmqcl",
      "type" : "function",
      "function" : {
        "name" : "getWeather",
        "arguments" : "{\"city\":\"Katowice\"}"
      }
    } ]
  }, {
    "role" : "tool",
    "content" : "Weather in Katowice: 22C, sunny",
    "tool_call_id" : "call_EIrKGD9Ec5BB3Ljw94tPmqcl"
  } ],
  "tools" : [ {
    "type" : "function",
    "function" : {
      "name" : "getCurrentTime",
      "description" : "Returns the current date and time",
      "parameters" : {
        "type" : "object",
        "properties" : { },
        "required" : [ ]
      }
    }
  }, {
    "type" : "function",
    "function" : {
      "name" : "calculateSquare",
      "description" : "Calculates the square of a given number",
      "parameters" : {
        "type" : "object",
        "properties" : {
          "number" : {
            "type" : "integer",
            "description" : "The number to square"
          }
        },
        "required" : [ "number" ]
      }
    }
  }, {
    "type" : "function",
    "function" : {
      "name" : "getWeather",
      "description" : "Gets mock weather for a city",
      "parameters" : {
        "type" : "object",
        "properties" : {
          "city" : {
            "type" : "string",
            "description" : "Name of the city"
          }
        },
        "required" : [ "city" ]
      }
    }
  } ],
  "temperature" : 0.3,
  "stream" : false
}
```

The request body as JSON

MESSAGES []

SYSTEM → written by dev, before the conversation

USER → written by the user

ASSISTANT → written by the model

TOOL → written by your app

ASSISTANT → written by the model

LOOP

SYSTEM ROLE

The system prompt is the job description.
Written once by a human, resent on every iteration.

```
{
  "role": "system",
  "content": "You are a helpful assistant. Use available tools when appropriate."
}
```

USER ROLE

The question.
Written by the user, resent on every iteration.

```
{
  "role": "user",
  "content": "How is the weather in Katowice?"
}
```

ASSISTANT ROLE

The model either asks for a tool, or gives the final answer.
A tool request means the loop continues.

finish_reason: "tool_calls"

```
{
  "role": "assistant",
  "content": null,
  "tool_calls": [{
    "id": "call_EIrKGD9E",
    "type": "function",
    "function": {
      "name": "getWeather",
      "arguments": "{\"city\":\"Katowice\"}"
    }
  }]
}
```

finish_reason: "stop"

```
{
  "role": "assistant",
  "content": "It's sunny in Katowice, 22°C."
}
```

TOOL ROLE

The result. Written by your app, never by the model.
`tool_call_id` matches the request, because the model may ask for several at once.

```
{
  "role": "tool",
  "content": "Weather in Katowice: 22C, sunny",
  "tool_call_id": "call_EIrKGD9E"
}
```

TOOLS

The functions the model may ask you to call.
Name, description and JSON schema that's all it knows about your code.

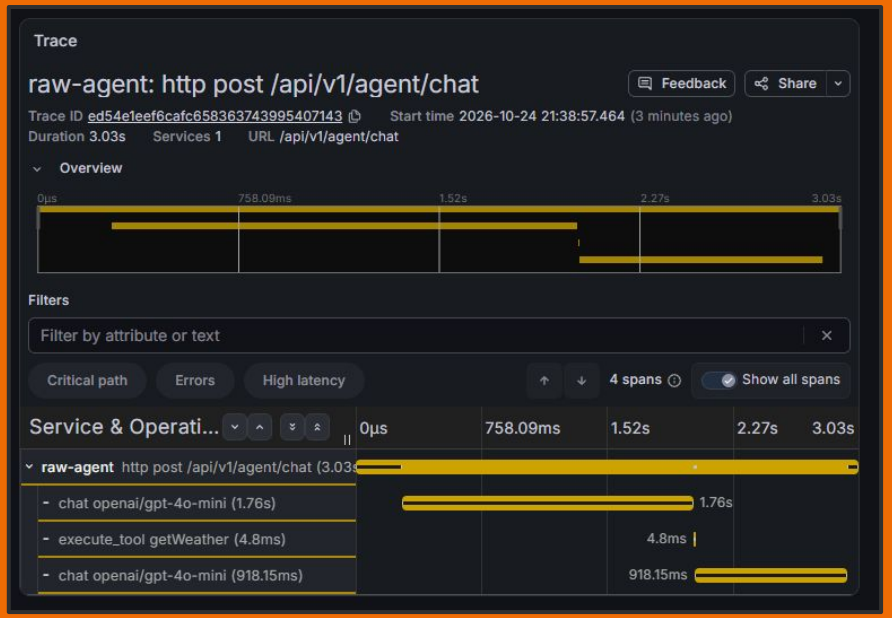
```json
{
  "type": "function",
  "function": {
    "name": "getWeather",
    "description": "Gets mock weather for a city",
    "parameters": {
      "type": "object",
      "properties": {
        "city": {
          "type": "string",
          "description": "Name of the city"
        }
      },
      "required": ["city"]
    }
  }
}
```

```java
package com.xkondix.rawagent.model;

import com.fasterxml.jackson.annotation.JsonInclude;
import tools.jackson.databind.node.ObjectNode;

@JsonInclude(JsonInclude.Include.NON_NULL)
public record ToolDefinition(
        String type,
        FunctionDef function
) {
    public record FunctionDef(
            String name,
            String description,
            ObjectNode parameters
    ) {}
}
```

Observability




```
2026-10-24 21:39:00.417  DEBUG  [LLM] Response status=200
2026-10-24 21:39:00.417  INFO   [LOOP] Done after 2 iteration(s)
2026-10-24 21:39:00.417  INFO   [LOOP] Tokens: input=150 output=15 total=165
2026-10-24 21:38:59.504  DEBUG  [LLM] POST https://openrouter.ai/api/v1/chat/completions
2026-10-24 21:38:59.504  DEBUG  [LLM] OpenAI request: model=openai/gpt-4o-mini messages=
  4 tools=3
2026-10-24 21:38:59.504  INFO   [LOOP] Iteration 2 — 4 messages in history
2026-10-24 21:38:59.504  INFO   [LOOP] Tool result: Weather in Katowice: 22C, sunny
2026-10-24 21:38:59.501  INFO   [LOOP] Tokens: input=115 output=16 total=131
2026-10-24 21:38:59.501  INFO   [LOOP] Tool call: getWeather args={"city":"Katowice"}
2026-10-24 21:38:59.501  INFO   [TOOL] Executing: getWeather args={"city":"Katowice"}
2026-10-24 21:38:59.444  DEBUG  [LLM] Response status=200
2026-10-24 21:38:57.770  DEBUG  [LLM] POST https://openrouter.ai/api/v1/chat/completions
2026-10-24 21:38:57.743  DEBUG  [LLM] OpenAI request: model=openai/gpt-4o-mini messages=
  2 tools=3
2026-10-24 21:38:57.735  INFO   [LOOP] Starting. User: How is the weather in Katowice?
2026-10-24 21:38:57.735  INFO   [LOOP] Iteration 1 — 2 messages in history
```

LangChain4j Agent

Everything AiServices does, LangChain4j also lets you do yourself.

Same library, same tools and same model.

Only the level of abstraction moves.



LangChain4j Abstraction

ChatModel	→	you run the loop
AiServices	→	the framework runs it
AgenticServices	→	you declare the workflow [part Ⅱ]

LangChain4j Loop

with AiServices

```
this.assistant = AiServices.builder(Assistant.class)
        .chatModel(model)
        .toolProvider(toolProvider)
        .maxToolCallingRoundTrips(10) // by default 100
        .chatMemoryProvider(memoryId ->
                MessageWindowChatMemory.withMaxMessages(20))
        .build();
```

`.toolProvider(toolProvider)`
→ a custom provider, wrapped for observability

`.tools(demoTools)`
→ simpler with no tool spans unless you hook in

without AiServices

```
while (iteration++ < maxIterations) {
    log.info("[RAW LOOP] Iteration {} — {} messages", iteration, messages.size());

    ChatResponse response = chatModel.chat(
            ChatRequest.builder()
                    .messages(messages)
                    .parameters(ChatRequestParameters.builder()
                            .toolSpecifications(toolExecutor.getToolSpecs())
                            .build())
                    .build());

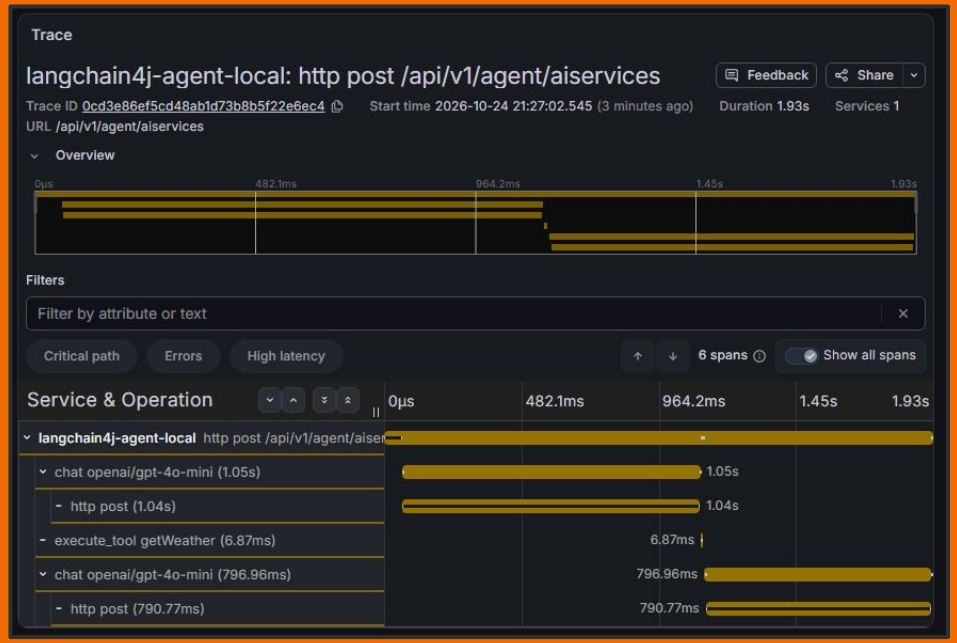
    AiMessage aiMessage = response.aiMessage();
    messages.add(aiMessage);

    if (!aiMessage.hasToolExecutionRequests()) {
        log.info("[RAW LOOP] Done after {} iterations", iteration);
        return aiMessage.text();
    }

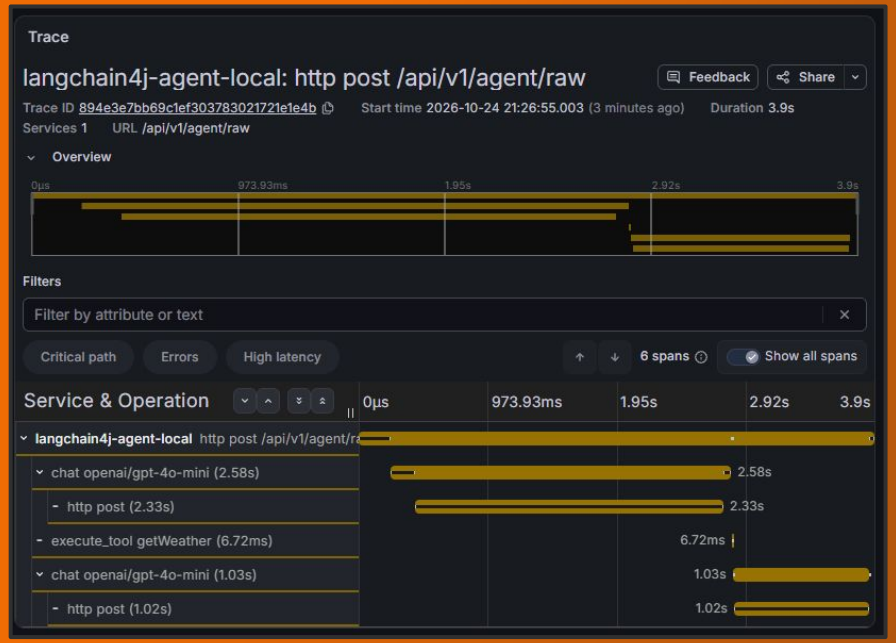
    for (var toolRequest : aiMessage.toolExecutionRequests()) {
        String result = executeToolWithSpan(
                toolRequest.name(), toolRequest.arguments(), iteration);
        messages.add(ToolExecutionResultMessage.from(toolRequest, result));
    }
}
```

Traces

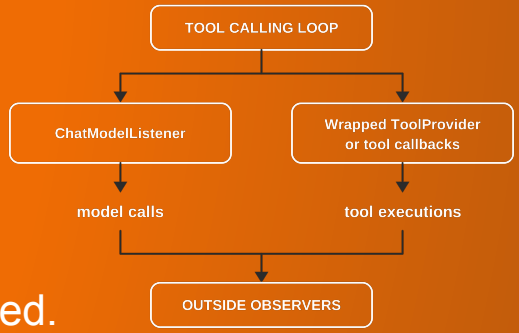
with AiServices



without AiServices



Logs



The loop is still observable, just from the outside.

A `ChatModelListener` sees every model call, tool requests included.

A wrapped `ToolProvider` or the `tool callbacks` see every tool execution.

with AiServices

```
: 2026-10-03 22:49:07.980  INFO  Tool executed: getWeather(Katowice)
: 2026-10-03 22:49:06.904  INFO  AiServices chat: userId=user-8gvin9
```

without AiServices

```
: 2026-10-03 22:44:42.018  INFO  [RAW LOOP] Done after 2 iterations
: 2026-10-03 22:44:40.916  INFO  [RAW LOOP] Iteration 2 — 4 messages
: 2026-10-03 22:44:40.914  INFO  [RAW LOOP] Result: Weather in Katowice: 20C, sunny
: 2026-10-03 22:44:40.913  INFO  Tool executed: getWeather(Katowice)
: 2026-10-03 22:44:40.911  INFO  [RAW LOOP] Tool: getWeather args={"city":"Katowice"}
: 2026-10-03 22:44:39.520  INFO  [RAW LOOP] Iteration 1 — 2 messages
```

Five ways around one loop

Objects && Functions	Scope	Can change the request or result?
ChatModelListener	every model call: request, response, error	no
chatRequestTransformer	every request, before it is sent	yes
ToolProvider wrapper	every tool call, around the execution	yes
beforeToolExecution / afterToolExecution	every tool call: request, then result, duration, failure	no
AI Service listeners	the whole call, as events	no

Two callbacks instead of a wrapper

Enough for logs and metrics.

Fragile for spans, because after never runs when the arguments don't parse.

The wrapper starts and ends the span in one try/finally.

```
Assistant assistant = AiServices.builder(Assistant.class)
        .chatModel(model)
        .tools(demoTools)
        .beforeToolExecution(before ->
                log.info("[TOOL] → {}", before.request().name()))
        .afterToolExecution(exec ->
                log.info("[TOOL] ← {} in {} ms, failed={}",
                        exec.request().name(),
                        exec.duration().toMillis(),
                        exec.hasFailed()))
        .build();
```

AgenticServices

More in chapter Ⅱ

```
var translator = AgenticServices.agentBuilder(Translator.class)
        .chatModel(chatModel)
        .maxToolCallingRoundTrips(10)
        .outputKey("translation")
        .build();

this.workflow = AgenticServices.sequenceBuilder(ScoutReportWorkflow.class)
        .subAgents(analyst, condenser, translator)
        .outputKey("translation")
        .build();
```

Spring AI

Spring AI doesn't hide the loop, it puts it in a chain you can reach

Same chain, same client, same model.

Only one advisor is the loop.

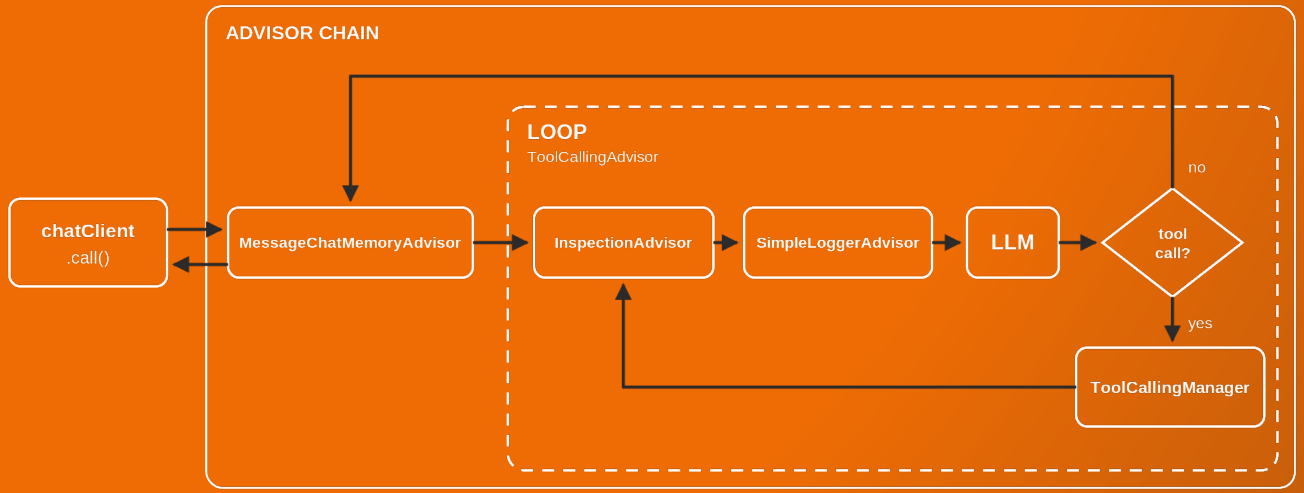


What is an Advisor?

An Advisor is middleware between your application and the AI model.

It works similarly to an aspect, allowing you to add behavior around AI interactions.

It can modify requests, process responses, and control the flow of interactions with the model, keeping this logic separate from your application code.



Advisor Code Snippet

```
public String chatWithAdvisors(String conversationId, String message) {
    log.info("Spring AI chat [advisors]: conversationId={}", conversationId);

    return chatClient.prompt()
            .user(message)
            .advisors(
                    new InspectionAdvisor(),
                    new SimpleLoggerAdvisor(),
                    MessageChatMemoryAdvisor.builder(chatMemory).build())
            .advisors(a -> a.param(ChatMemory.CONVERSATION_ID, conversationId))
            .tools(demoFunctions)
            .call()
            .content();
}
```

Adding advisors to the chain

Parameter configuration for advisors (linking an identifier with a message)

Before and After

```java
@Override
public int getOrder() {
    return 0;
}

@Override
public ChatClientRequest before(
        ChatClientRequest request,
        AdvisorChain advisorChain) {

    log.info("[Advisor] BEFORE – {} messages, user=\"{}\"",
            request.prompt().getInstructions().size(),
            request.prompt().getUserMessage().getText());

    return request;
}

@Override
public ChatClientResponse after(
        ChatClientResponse response,
        AdvisorChain advisorChain) {

    var result = response.chatResponse().getResult();
    String text = result.getOutput().getText();
    boolean asksForTool = result.getOutput().hasToolCalls();

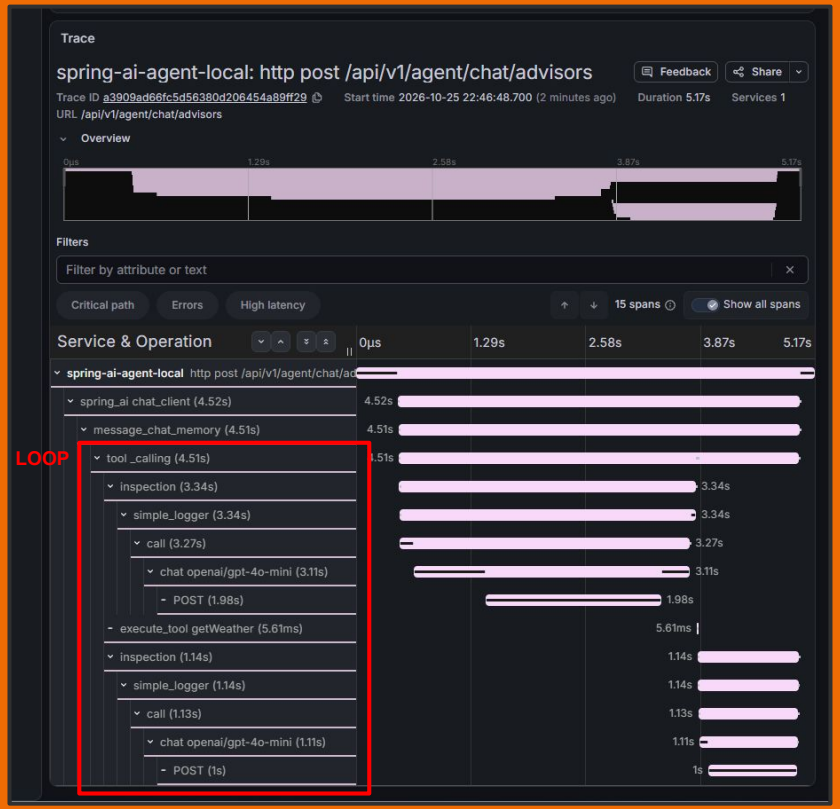
    log.info("[Advisor] AFTER – toolCalls={}, responseLength={}",
            asksForTool, text != null ? text.length() : 0);

    return response;
}
```

```
ected_level scope_name   Show original line

: 2026-10-04 18:46:14.824  INFO  info  com.xkondix.springai.agent.advisor.InspectionAdvisor
: 2026-10-04 18:46:14.822  DEBUG  debug  org.springframework.ai.chat.client.advisor.SimpleLoggerAdvisor
: 2026-10-04 18:46:14.820  INFO  info  org.springframework.ai.chat.observation.ChatModelCompletionObservat
ionHandler
: 2026-10-04 18:46:14.820  INFO  info  org.springframework.ai.chat.observation.ChatModelPromptContentObser
vationHandler
: 2026-10-04 18:46:13.053  DEBUG  debug  org.springframework.ai.chat.client.advisor.SimpleLoggerAdvisor
: 2026-10-04 18:46:13.050  INFO  info  com.xkondix.springai.agent.advisor.InspectionAdvisor
: 2026-10-04 18:46:13.046  INFO  info  com.xkondix.springai.agent.tools.DemoFunctions
: 2026-10-04 18:46:13.036  INFO  info  com.xkondix.springai.agent.advisor.InspectionAdvisor  1
: 2026-10-04 18:46:13.035  DEBUG  debug  org.springframework.ai.chat.client.advisor.SimpleLoggerAdvisor  2
: 2026-10-04 18:46:12.983  INFO  info  org.springframework.ai.chat.observation.ChatModelCompletionObservat
ionHandler
: 2026-10-04 18:46:12.983  INFO  info  org.springframework.ai.chat.observation.ChatModelPromptContentObser
vationHandler
: 2026-10-04 18:46:10.480  DEBUG  debug  org.springframework.ai.chat.client.advisor.SimpleLoggerAdvisor  2
: 2026-10-04 18:46:10.477  INFO  info  com.xkondix.springai.agent.advisor.InspectionAdvisor  1
: 2026-10-04 18:46:10.343  INFO  info  com.xkondix.springai.agent.service.SpringAiAgentService
```

Advisors Inside the Loop




```
sion=1.55.0 trace_id=97c9a5af90cc954c9158b4…

: 2026-10-04 15:43:22.369  INFO  [Advisor] AFTER – toolCalls=false, responseLength=52
: 2026-10-04 15:43:22.369  INFO  Chat Model Completion:
  ["The weather in Katowice is currently 20°C and sunny."]
: 2026-10-04 15:43:22.369  INFO  Chat Model Prompt Content:
  ["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowi
  ce?", "", ""]
: 2026-10-04 15:43:20.734  INFO  [Advisor] BEFORE – 4 messages, user="How is the weather in Katow
  ice?"
: 2026-10-04 15:43:20.732  INFO  Spring AI Tool: getWeather(Katowice)
: 2026-10-04 15:43:20.728  INFO  [Advisor] AFTER – toolCalls=true, responseLength=0
: 2026-10-04 15:43:20.718  INFO  Chat Model Completion:
  []
: 2026-10-04 15:43:20.718  INFO  Chat Model Prompt Content:
  ["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowi
  ce?"]
: 2026-10-04 15:43:19.144  INFO  [Advisor] BEFORE – 2 messages, user="How is the weather in Katow
  ice?"
: 2026-10-04 15:43:19.129  INFO  Spring AI chat [advisors]: conversationId=user-f9lvbd
```

Advisor Order

The execution order of advisors in the chain is determined by the `getOrder()` method.

Lower order = earlier in the chain, farther from the model [LLM].

Higher order = later in the chain, closer to the model [LLM].



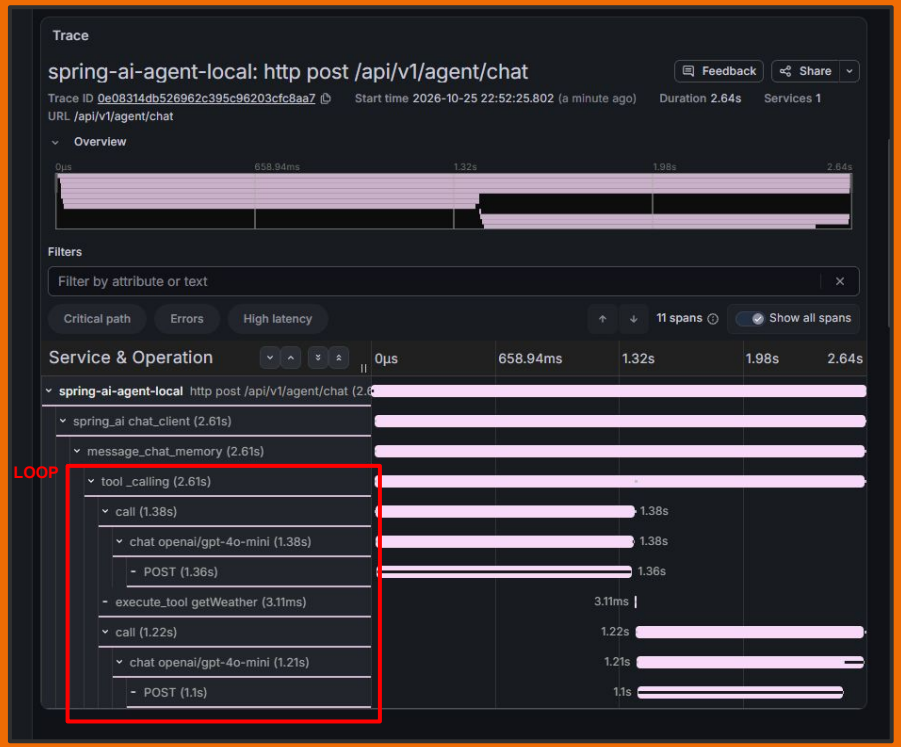
Example of Advisors

	Order ID Value
MessageChatMemoryAdvisor → adds history as separate messages	**MIN_VALUE + 200**
QuestionAnswerAdvisor → adds context from a vector store (RAG)	**0**
SafeGuardAdvisor → checks the prompt and history for banned words	**0**
ToolCallingAdvisor → owns the tool calling lifecycle	**MIN_VALUE + 300**
SimpleLoggerAdvisor → logs request and response information	**0**

Integer.MIN_VALUE = -2 147 483 648 ← HIGHEST_PRECEDENCE (doc name)

Order 0 is not first. It is inside the loop!!!

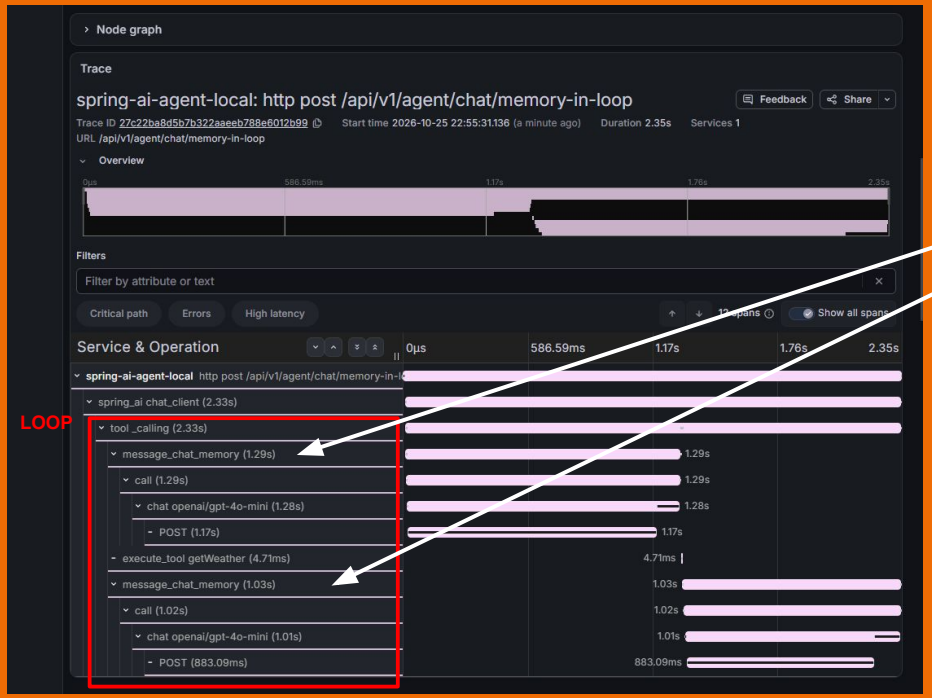
Default



One memory advisor, outside the loop. It saves only the question and the final answer.

```
2026-10-25 22:52:28.426  INFO  Chat Model Prompt Content:
["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowice?", "", ""]
2026-10-25 22:52:28.423  INFO  Chat Model Completion:
["The weather in Katowice is currently 20°C and sunny."]
2026-10-25 22:52:27.200  INFO  Spring AI Tool: getWeather(Katowice)
2026-10-25 22:52:27.200  INFO  Chat Model Completion:
[]
2026-10-25 22:52:27.200  INFO  Chat Model Prompt Content:
["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowice?"]
2026-10-25 22:52:25.804  INFO  Spring AI chat [default]: conversationId=user-3xh2yb
```

Memory in the Loop



Memory inside the loop stores the whole tool transcript.

```
: 2026-10-25 22:55:33.476  INFO  Chat Model Prompt Content:
["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowice?", "", ""]
: 2026-10-25 22:55:33.475  INFO  Chat Model Completion:
["The weather in Katowice is currently 20°C and sunny."]
: 2026-10-25 22:55:32.441  INFO  Spring AI Tool: getWeather(Katowice)
: 2026-10-25 22:55:32.433  INFO  Chat Model Prompt Content:
["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowice?"]
: 2026-10-25 22:55:32.430  INFO  Chat Model Completion:
[]
: 2026-10-25 22:55:31.140  INFO  Spring AI chat [memory in loop]: conversationId=user-mwixii
```

Memory Inside && Outside

Now one memory advisor is outside the loop and one inside.
Both save the question, so from the second iteration on it is sent twice.
Spring AI catches duplicates between the loop's history and memory, but not between two memory advisors.

```java
@Bean
public ChatClient chatClient(ChatClient.Builder builder, MessageWindowChatMemory chatMemory) {
    return builder
            .defaultSystem("You are a helpful assistant. " +
                    "Use available tools when needed.")
            .defaultAdvisors(
                    MessageChatMemoryAdvisor.builder(chatMemory).build())
            .build();
}
```

```java
public String chatWithMemoryInLoop(String conversationId, String message) {
    log.info("Spring AI chat [memory in loop]: conversationId={}", conversationId);

    return chatClient.prompt()
            .user(message)
            .advisors(MessageChatMemoryAdvisor.builder(chatMemory)
                    .order(MEMORY_INSIDE_LOOP)
                    .build())
            .advisors(a -> a.param(ChatMemory.CONVERSATION_ID, conversationId))
            .tools(demoFunctions)
            .call()
            .content();
}
```

```
:d0bbb3d04f0166d6a…
: 2026-10-04 15:43:53.762  INFO  Chat Model Completion:
["The weather in Katowice is currently 20°C and sunny."]
: 2026-10-04 15:43:53.762  INFO  Chat Model Prompt Content:
["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowice?", "How is the weather in Katowice?", "", ""]
: 2026-10-04 15:43:52.920  INFO  Spring AI Tool: getWeather(Katowice)
: 2026-10-04 15:43:52.916  INFO  Chat Model Prompt Content:
["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowice?"]
: 2026-10-04 15:43:52.916  INFO  Chat Model Completion:
[]
: 2026-10-04 15:43:51.870  INFO  Spring AI chat [memory in loop]: conversationId=user-gco646
```

Three levels of hiding the loop

BY HAND

raw-agent

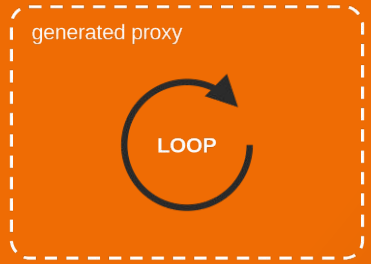
```
while (i++ < MAX) {
  resp = model.call(messages);
  if (!resp.hasToolCalls()) return resp;
  for (call : resp.toolCalls())
    messages.add(execute(call));
}
```



FRAMEWORK SERVICE

LangChain4j AiServices

```
interface Assistant {
  String chat(String msg);
}
AiServices.builder(Assistant.class)
  .chatModel(model).tools(tools)
  .build();
```



ADVISOR CHAIN

Spring AI ChatClient

```
chatClient.prompt(msg)
  .tools(tools)
  .call()
  .content();
```



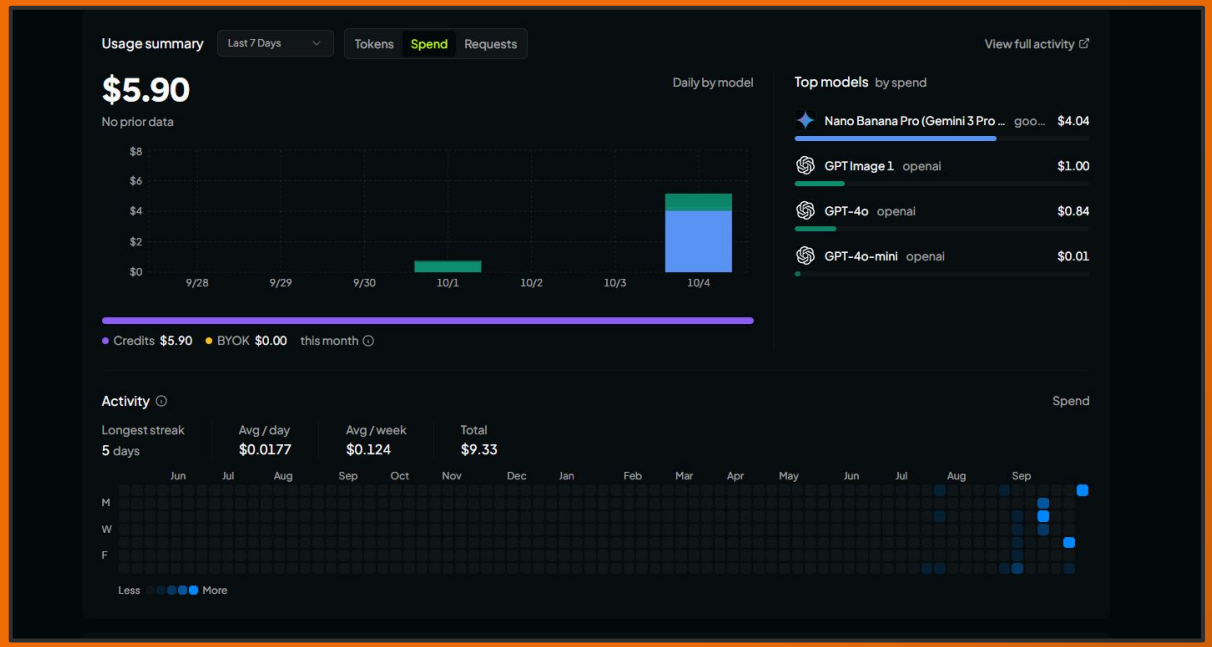
Who stops the loop?

Implementation	Limit	When exceeded
raw-agent	10 iterations	a string in HTTP 200, indistinguishable from an answer
LC4j AiServices	100 by default (10 in the demo)	exception
Spring AI 2.0.1	40 per tool / 150 per request **doc**	response with a dedicated finish reason

OpenRouter

OpenRouter is a single API in front of hundreds of models from dozens of vendors.

One key, one base URL, and the OpenAI protocol everyone already speaks.



Configuration

Spring AI configuration

Own properties

```yaml
spring:
  config:
    activate:
      on-profile: openrouter
  ai:
    model:
      chat: openai
    openai:
      base-url: https://openrouter.ai/api/v1
      api-key: ${OPENROUTER_API_KEY}
      chat:
        model: ${OPENROUTER_MODEL:openai/gpt-4o-mini}
        temperature: 0.3
      image:
        model: ${OPENROUTER_IMAGE_MODEL:google/gemini-3-pro-image}
        size: 1024x1024
      audio:
        speech:
          model: ${OPENROUTER_TTS_MODEL:x-ai/grok-voice-tts-1.0}
          voice: ${OPENROUTER_TTS_VOICE:eve}
          response-format: mp3
          speed: 1.0
        transcription:
          model: ${OPENROUTER_STT_MODEL:openai/whisper-large-v3}
```

```yaml
multimodal:
  media-dir: ./media
  vision-model: ${OPENROUTER_VISION_MODEL:openai/gpt-4o}
  music:
    model-pro: ${OPENROUTER_MUSIC_MODEL_PRO:google/lyria-3-pro-preview}
    model-clip: ${OPENROUTER_MUSIC_MODEL_CLIP:google/lyria-3-clip-preview}
  video:
    enabled: ${MULTIMODAL_VIDEO_ENABLED:false}
    model: ${OPENROUTER_VIDEO_MODEL:google/veo-3.1}
    resolution: ${OPENROUTER_VIDEO_RESOLUTION:720p}
```

Chat

```
public String chat(String conversationId, String question, String imageUrl) {
    String message = (imageUrl == null || imageUrl.isBlank())
            ? question
            : question + "\n\n[The user attached an image at " + imageUrl + "]";

    log.info("[MM] router conversationId={} image={} question='{}'",
            conversationId, imageUrl, question);

    List<Object> toolBeans = new ArrayList<>(List.of(tools, musicTools));
    if (videoTools != null) {
        toolBeans.add(videoTools);
    }

    return router.prompt()
            .advisors(MessageChatMemoryAdvisor.builder(chatMemory).build())
            .advisors(a -> a.param(ChatMemory.CONVERSATION_ID, conversationId))
            .tools(toolBeans.toArray())
            .user(message)
            .call()
            .content();
}
```

Image Generation

```
try {
    ImageResponse response = imageModel.call(new ImagePrompt(description));
    var image = response.getResult().getOutput();

    byte[] bytes;
    if (image.getB64Json() != null && !image.getB64Json().isBlank()) {
        bytes = Base64.getDecoder().decode(image.getB64Json());
    } else if (image.getUrl() != null && !image.getUrl().isBlank()) {
        bytes = download(image.getUrl());
    } else {
        collector.recordFailure();
        record("image", 0, sample);
        return "ERROR: the image model returned neither b64_json nor a url. Do NOT retry.";
    }
```

Speech

```
public String speak_text(
        @ToolParam(description = "The exact text to read out loud, in the language it is written in")
        String text) {

    String refusal = refuseIfBudgetSpent("speak_text", "read text out loud");
    if (refusal != null) {
        return refusal;
    }

    log.info("[MM] speak_text: {} chars", text == null ? 0 : text.length());
    Timer.Sample sample = Timer.start(registry);
    try {
        byte[] mp3 = speechModel.call(text);
        String url = storage.store(mp3, "mp3");
        String scriptUrl = storage.store(text.getBytes(StandardCharsets.UTF_8), "txt");
        collector.add(Artifact.of(scriptUrl, Artifact.Kind.TEXT, abbreviate(text)));
        collector.add(Artifact.of(url, Artifact.Kind.SPEECH, abbreviate(text)));

        record("speech", mp3.length, sample);
        return "Done — the audio has been generated and is playing for the user.";

    } catch (RuntimeException e) {
        int failures = collector.recordFailure();
        log.error("[MM] speak_text failed ({}/{}): {}",
                failures, MediaCollector.MAX_FAILURES, e.getMessage());
        record("speech", 0, sample);
        return "ERROR: could not synthesize speech — " + e.getMessage()
                + ". Do NOT retry; tell the user this part failed.";
    }
}
```

Vision

```
public String analyze(String imageUrl, String question) {
    String relative = imageUrl.replace(MediaStorage.URL_PREFIX, "");
    Path file = storage.resolve(relative);
    log.info("[MM] analyze_image {} question='{}'", relative, question);

    Timer.Sample sample = Timer.start(registry);
    try {
        MimeType mime = relative.endsWith(".png")
                ? MimeTypeUtils.IMAGE_PNG : MimeTypeUtils.IMAGE_JPEG;

        String answer = vision.prompt()
                .options(OpenAiChatOptions.builder().model(model))
                .user(u -> u
                        .text(question == null || question.isBlank()
                                ? "Describe this image."
                                : question)
                        .media(new Media(mime, new FileSystemResource(file))))
                .call()
                .content();

        record(sample);
        return answer;

    } catch (RuntimeException e) {
        log.error("[MM] analyze_image failed: {}", e.getMessage());
        record(sample);
        throw e;
    }
}
```

Music

Here is where the trouble started….

Lyria is a chat endpoint. Getting the bytes out took five attempts.

1. chatModel.call() → timed lyrics, no audio “[0.0:7.3] O-Zone, we dance tonight…]“
2. + getMedia() → media blocks = 0
3. + modalities → `voice` is required, but was not set
4. + voice → 400: Audio output requires stream: true
5. + stream() → the stream arrives, media blocks = 0

Every one of those errors was correct and specific. None of them said where to look.

Why can't Spring AI reach it?

Layer	Behavior
Google	Exposes Lyria through the Interactions API, not `/audio/speech`
OpenRouter	Returns audio from chat only when `stream: true`
OpenAI SDK	Validates a speech-oriented request schema and requires `voice`
Spring AI	Maps audio to `Media` only in the non-streaming (blocking) path

No layer made a mistake, so no layer could warn.
Music dropped to raw HTTP.
One modality on the transport is cheaper than pretending the abstraction covers five.

Code Snippet

```
public Result generate(String model, String prompt) {
    Map<String, Object> body = new LinkedHashMap<>();
    body.put("model", model);
    body.put("messages", List.of(Map.of("role", "user", "content", prompt)));
    body.put("modalities", List.of("text", "audio"));
    body.put("audio", Map.of("format", "mp3"));
    body.put("stream", true);

    HttpRequest request = HttpRequest.newBuilder(URI.create(baseUrl + "/chat/completions"))
            .header("Authorization", "Bearer " + apiKey)
            .header("Content-Type", "application/json")
            .header("Accept", "text/event-stream")
            .timeout(Duration.ofMinutes(5))
            .POST(HttpRequest.BodyPublishers.ofString(mapper.writeValueAsString(body)))
            .build();

    try {
        HttpResponse<Stream<String>> response =
                http.send(request, HttpResponse.BodyHandlers.ofLines());

        if (response.statusCode() != 200) {
            throw new IllegalStateException("Lyria returned HTTP " + response.statusCode());
        }
        return consume(response.body());
```

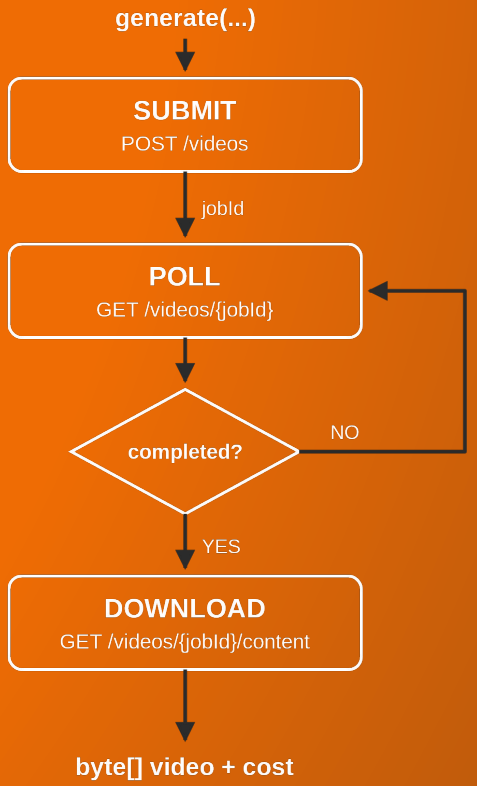
Bytes

Video

Video generation is asynchronous: submit, poll, download.

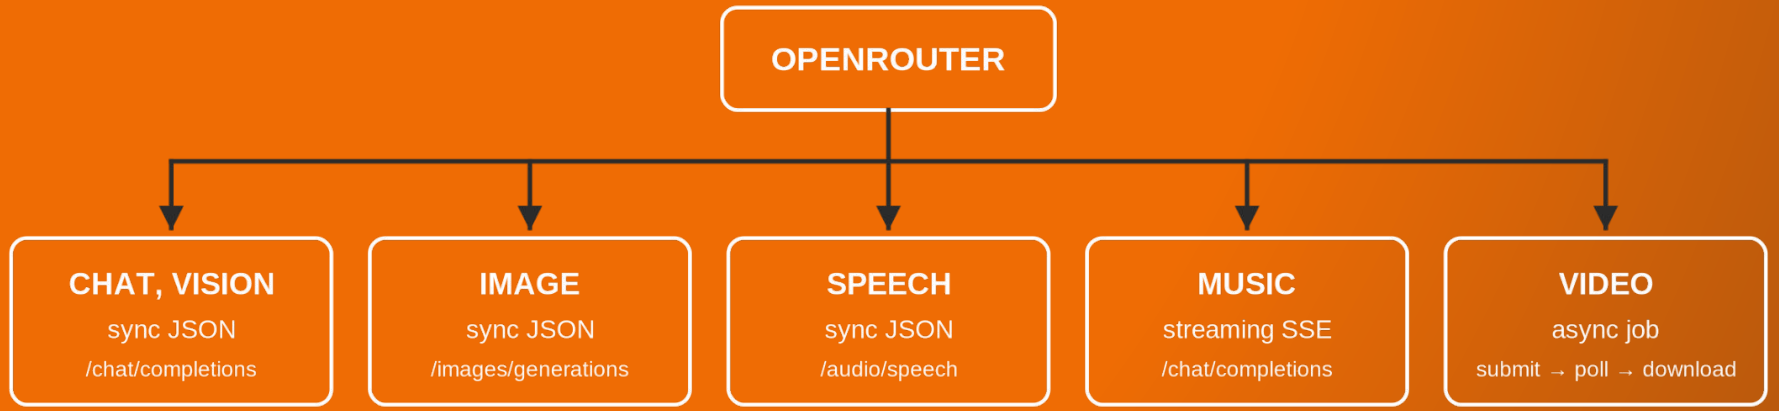
There is no `VideoModel` in Spring AI, so the client is custom.

Music had the opposite problem, the abstraction existed and could not be reached.



Is one protocol really enough?

One provider does not mean one protocol.



See you in part Ⅱ next week!

