


The Agent Is a Loop

Author: Konrad Kowalczyk

Repository: [ai-agents-JAVA-SPRING](#)

Let's start from the beginning

This is a continuation of last year's presentation, which focused on...



OLD REPO → [xkondix AI Sandbox](#)

Roadmap

1. The Agent Is a Loop

Writing one by hand, and what frameworks hide when they write it for you

2. Agentic Workflow Patterns

Five shapes, two frameworks, and when the DSL earns its keep

3. MCP as a Process Boundary

Tools, resources and prompts over two transports: HTTP for agents, STDIO for Claude Desktop

4. How Do You Know It Works?

Traces, metrics, logs and dashboards for everything your agents actually do


The Agent Is a Loop


1. **What is an agent?**
2. **Build it** -> No framework
3. **Hide it** -> LangChain4j
4. **Control it** -> Spring AI
5. **Mix it** -> One provider, ∞ protocols
6. **Pay for it** -> The cost tokens don't show



What is an agent?







Raw Agent

The loop exits when the model stops requesting tools.

stop condition = no more tool calls

sources  openai
strandsagents

```
// — Agent loop —  
for (int i = 1; i <= MAX_ITERATIONS; i++) {  
    log.info("[LOOP] Iteration {} - {} messages in history", i, history.size());  
  
    // — Step 2: Call LLM —  
    var response = llmClient.chat(history, tools.getDefinitions());  
    var message = response.firstMessage();  
  
    if (message == null) {  
        log.error("[LOOP] Empty response from LLM*");  
        return "Error: empty response from model.";  
    }  
  
    history.add(Message.assistant(message.content(), message.toolCalls()));  
  
    if (response.usage() != null) {  
        log.info("[LOOP] Tokens: input={} output={} total={}",  
            response.usage().promptTokens(),  
            response.usage().completionTokens(),  
            response.usage().totalTokens());  
    }  
  
    // — Step 3: Tool calls? —  
    if (!response.hasToolCalls()) {  
        log.info("[LOOP] Done after {} iteration(s)", i);  
        return message.content() != null ? message.content() : "";  
    }  
  
    for (ToolCall toolCall : message.toolCalls()) {  
        String result = executeToolWithSpan(toolCall, i);  
        history.add(Message.tool(toolCall.id(), result));  
    }  
}  
  
return "Max iterations (" + MAX_ITERATIONS + ") reached.";
```

The request body structure

Every exchange with the model is JSON. The framework only decides who writes it.

- MODEL** → the model that handles this request
- MESSAGES** → the whole conversation, resent on every call
- TOOLS** → your Java methods, exposed to the model
- TEMPERATURE** → how deterministic the answer is
- STREAM** → false: one complete response
true: chunks as they are generated

```
{  
  "model": "openai/gpt-4o-mini",  
  "messages": [ ... ],  
  "tools": [ ... ],  
  "temperature": 0.3,  
  "stream": false  
}
```

The request body as JSON

MESSAGES []

SYSTEM → the instructions

USER → the question

ASSISTANT → a request for a tool

TOOL → the result your code produced

ASSISTANT → the answer

LOOP

```
{
  "model": "openai/gpt-4o-mini",
  "messages": [ {
    "role": "system",
    "content": "You are a helpful assistant. Use available tools when appropriate."
  }, {
    "role": "user",
    "content": "How is the weather in Katowice?"
  }, {
    "role": "assistant",
    "tool_calls": [ {
      "id": "call_EIrK029Ec5883Ljw94tPmqcl",
      "type": "function",
      "function": {
        "name": "getWeather",
        "arguments": "{ \"city\": \"Katowice\"}"
      }
    }
  ]
}, {
  "role": "tool",
  "content": "Weather in Katowice: 22C, sunny",
  "tool_call_id": "call_EIrK029Ec5883Ljw94tPmqcl"
} ],
  "tools": [ {
    "type": "function",
    "function": {
      "name": "getCurrentTime",
      "description": "Returns the current date and time",
      "parameters": {
        "type": "object",
        "properties": { },
        "required": [ ]
      }
    }
  }, {
    "type": "function",
    "function": {
      "name": "calculateSquare",
      "description": "Calculates the square of a given number",
      "parameters": {
        "type": "object",
        "properties": {
          "number": {
            "type": "integer",
            "description": "The number to square"
          }
        },
        "required": [ "number" ]
      }
    }
  }, {
    "type": "function",
    "function": {
      "name": "getWeather",
      "description": "Gets mock weather for a city",
      "parameters": {
        "type": "object",
        "properties": {
          "city": {
            "type": "string",
            "description": "Name of the city"
          }
        },
        "required": [ "city" ]
      }
    }
  ]
},
  "temperature": 0.3,
  "stream": false
}
```

The request body as JSON

MESSAGES []

SYSTEM → written by dev, before the conversation

USER → written by the user

ASSISTANT → written by the model

TOOL → written by your the app

ASSISTANT → written by the model

┌

└

LOOP

SYSTEM ROLE

The system prompt is the job description.
Written once, before anyone asks anything.

```
{  
  "role": "system",  
  "content": "You are a helpful assistant. Use available tools when appropriate."  
}
```

USER ROLE

The question.

The only message a human wrote. It never appears again in the loop.

```
{  
  "role": "user",  
  "content": "How is the weather in Katowice?"  
}
```

ASSISTANT ROLE

The model either asks for a tool, or gives the final answer.
A tool request means the loop continues.

```
{
  "role": "assistant",
  "tool_calls": [{
    "id": "call_EIrKGD9E",
    "function": {
      "name": "getWeather",
      "arguments": "{\"city\":\"Katowice\"}"
    }
  }]
}
```

TOOL ROLE

The result. Written by your app, never by the model.

tool_call_id matches the request, because the model may ask for several at once.

```
{  
  "role": "tool",  
  "content": "Weather in Katowice: 22C, sunny",  
  "tool_call_id": "call_EIrKGD9E"  
}
```

TOOLS

The list of functions that the model may call, each with a description.
The description is the only thing the model knows about your code.


```
{
  "type": "function",
  "function": {
    "name": "getWeather",
    "description": "Gets mock weather for a city",
    "parameters": {
      "type": "object",
      "properties": {
        "city": {
          "type": "string",
          "description": "Name of the city"
        }
      },
      "required": ["city"]
    }
  }
}
```

```
package com.xkondix.rawagent.model;

import com.fasterxml.jackson.annotation.JsonInclude;
import tools.jackson.databind.node.ObjectNode;

19 usages  ⚡ Konrad
@JsonInclude(JsonInclude.Include.NON_NULL)
public record ToolDefinition(
    no usages
    String type,
    3 usages
    FunctionDef function
) {
    4 usages  ⚡ Konrad
    public record FunctionDef(
        2 usages
        String name,
        no usages
        String description,
        1 usage
        ObjectNode parameters
    ) {}
}
```

Observability



```
: 2026-10-04 00:16:26.131 DEBUG [LLM] Response status=200
: 2026-10-04 00:16:26.131 INFO [LOOP] Done after 2 iteration(s)
: 2026-10-04 00:16:26.131 INFO [LOOP] Tokens: input=150 output=15 total=165
: 2026-10-04 00:16:25.114 DEBUG [LLM] POST https://openrouter.ai/api/v1/chat/completions
: 2026-10-04 00:16:25.103 DEBUG [LLM] OpenAI request: model=openai/gpt-4o-mini messages=4 tools=3
: 2026-10-04 00:16:25.103 INFO [LOOP] Tool result: Weather in Katowice: 22C, sunny
: 2026-10-04 00:16:25.103 INFO [LOOP] Tool call: getWeather args={"city":"Katowice"}
: 2026-10-04 00:16:25.103 INFO [LOOP] Iteration 2 - 4 messages in history
: 2026-10-04 00:16:25.103 INFO [LOOP] Tokens: input=115 output=16 total=131
: 2026-10-04 00:16:25.103 INFO [TOOL] Executing: getWeather args={"city":"Katowice"}
: 2026-10-04 00:16:25.082 DEBUG [LLM] Response status=200
: 2026-10-04 00:16:23.148 DEBUG [LLM] POST https://openrouter.ai/api/v1/chat/completions
: 2026-10-04 00:16:23.132 DEBUG [LLM] OpenAI request: model=openai/gpt-4o-mini messages=2 tools=3
: 2026-10-04 00:16:23.132 INFO [LOOP] Iteration 1 - 2 messages in history
: 2026-10-04 00:16:23.132 INFO [LOOP] Starting. User: How is the weather in Katowice?
```

LangChain4j Agent

Everything AiServices does, LangChain4j also lets you do yourself.

Same library, same tools, same model. Only the level of abstraction moves.



LangChain4j Abstraction

ChatModel → you run the loop

AiServices → the framework runs it

AgenticServices → you declare where it stops [part II]

LangChain4j Loop

with AiServices

```
this.assistant = AiServices.builder(Assistant.class)
    .chatModel(model)
    .toolProvider(toolProvider)
    .maxToolCallingRoundTrips(10) //by default 100
    .chatMemoryProvider(memoryId ->
        MessageWindowChatMemory.withMaxMessages(20))
    .build();
```

`.toolProvider(toolProvider)` is a custom provider,
wrapped for observability

`.tools(demoTools)` is simpler, but
invisible in traces

without AiServices

```
while (iteration++ < maxIterations) {
    log.info("[RAW LOOP] Iteration {} - {} messages", iteration, messages.size());

    ChatResponse response = chatModel.chat(
        ChatRequest.builder()
            .messages(messages)
            .parameters(ChatRequestParameters.builder() Builder<capture of ?>
                .toolSpecifications(toolExecutor.getToolSpecs()) capture of ?
                .build())
            .build());


    AiMessage aiMessage = response.aiMessage();
    messages.add(aiMessage);

    if (!aiMessage.hasToolExecutionRequests()) {
        log.info("[RAW LOOP] Done after {} iterations", iteration);
        return aiMessage.text();
    }


    for (var toolRequest : aiMessage.toolExecutionRequests()) {
        String result = executeToolWithSpan(
            toolRequest.name(), toolRequest.arguments(), iteration);
        messages.add(ToolExecutionResultMessage.from(toolRequest, result));
    }
}
```

Traces

with AiServices



without AiServices



Logs

The loop is still observable, just not from the inside.
You can register a listener; you don't add a line.

with AiServices

```
: 2026-10-03 22:49:07.980 INFO Tool executed: getWeather(Katowice)
: 2026-10-03 22:49:06.904 INFO AiServices chat: userId=user-8gvin9
```

without AiServices

```
: 2026-10-03 22:44:42.018 INFO [RAW LOOP] Done after 2 iterations
: 2026-10-03 22:44:40.916 INFO [RAW LOOP] Iteration 2 - 4 messages
: 2026-10-03 22:44:40.914 INFO [RAW LOOP] Result: Weather in Katowice: 20C, sunny
: 2026-10-03 22:44:40.913 INFO Tool executed: getWeather(Katowice)
: 2026-10-03 22:44:40.911 INFO [RAW LOOP] Tool: getWeather args={"city":"Katowice"}
: 2026-10-03 22:44:39.520 INFO [RAW LOOP] Iteration 1 - 2 messages
```

AgenticServices

More in chapter II

```
var translator = AgenticServices.agentBuilder(Translator.class) AgentBuilder<Translator, AgentBuilder<...>>
    .chatModel(chatModel) AgentBuilder<Translator, capture of ?>
    .maxToolCallingRoundTrips(10) capture of ?
    .outputKey("translation")
    .build();

this.workflow = AgenticServices.sequenceBuilder(ScoutReportWorkflow.class)
    .subAgents(analyst, condenser, translator)
    .outputKey(s: "translation")
    .build();
```

Spring AI

Everything the loop does, Spring AI lets you reach.


Same chain, same client, same model. Only one advisor is the loop.



What an Advisor is?

An Advisor is middleware between your application and the AI model. It works similarly to an aspect, allowing you to add behavior around AI interactions.

It can modify requests, process responses, and control the flow of interactions with the model, keeping this logic separate from your application code.



Advisor Code Snippet

```
public String chatWithAdvisors(String conversationId, String message) {  
    log.info("Spring AI chat [advisors]: conversationId={}", conversationId);  
  
    return chatClient.prompt() ChatClientRequestSpec  
        .user(message)  
        .advisors(  
            new InspectionAdvisor(),  
            new SimpleLoggerAdvisor(),  
            MessageChatMemoryAdvisor.builder(chatMemory).build()  
        ).advisors(a -> a.param(ChatMemory.CONVERSATION_ID, conversationId))  
        .tools(demoFunctions)  
        .call() CallResponseSpec  
        .content();  
}
```

Adding advisors to the chain

Parameter configuration for advisors (linking an identifier with a message)

Before and After

```
@Override
public int getOrder() {
    return 0;
}

@Override
public ChatClientRequest before(
    ChatClientRequest request,
    AdvisorChain advisorChain) {

    log.info("[Advisor] BEFORE - {} messages, user=\"{}\"",
        request.prompt().getInstructions().size(),
        request.prompt().getUserMessage().getText());

    return request;
}

@Override
public ChatClientResponse after(
    ChatClientResponse response,
    AdvisorChain advisorChain) {


    var result = response.chatResponse().getResult();
    String text = result.getOutput().getText();
    boolean asksForTool = result.getOutput().hasToolCalls();

    log.info("[Advisor] AFTER - toolCalls={}, responseLength={}",
        asksForTool, text != null ? text.length() : 0);

    return response;
}
```

```
ected_level scope_name Show original line
: 2026-10-04 18:46:14.824 INFO info com.xkondix.springai.agent.advisor.InspectionAdvisor
: 2026-10-04 18:46:14.822 DEBUG debug org.springframework.ai.chat.client.advisor.SimpleLoggerAdvisor
: 2026-10-04 18:46:14.820 INFO info org.springframework.ai.chat.observation.ChatModelCompletionObservat
ionHandler
: 2026-10-04 18:46:14.820 INFO info org.springframework.ai.chat.observation.ChatModelPromptContentObser
vationHandler
: 2026-10-04 18:46:13.053 DEBUG debug org.springframework.ai.chat.client.advisor.SimpleLoggerAdvisor
: 2026-10-04 18:46:13.050 INFO info com.xkondix.springai.agent.advisor.InspectionAdvisor
: 2026-10-04 18:46:13.046 INFO info com.xkondix.springai.agent.tools.DemoFunctions
: 2026-10-04 18:46:13.036 INFO info com.xkondix.springai.agent.advisor.InspectionAdvisor 1
: 2026-10-04 18:46:13.035 DEBUG debug org.springframework.ai.chat.client.advisor.SimpleLoggerAdvisor 2
: 2026-10-04 18:46:12.983 INFO info org.springframework.ai.chat.observation.ChatModelCompletionObservat
ionHandler
: 2026-10-04 18:46:12.983 INFO info org.springframework.ai.chat.observation.ChatModelPromptContentObser
vationHandler
: 2026-10-04 18:46:10.480 DEBUG debug org.springframework.ai.chat.client.advisor.SimpleLoggerAdvisor 2
: 2026-10-04 18:46:10.477 INFO info com.xkondix.springai.agent.advisor.InspectionAdvisor 1
: 2026-10-04 18:46:10.343 INFO info com.xkondix.springai.agent.service.SpringAiAgentService
```

Advisor (debug logging off)



```
size=1.55.0 trace_id=97c9a5af90cc954c9158b4481bcc4ae1
- : 2026-10-04 15:43:22.369 INFO [Advisor] AFTER - toolCalls=false, responseLength=52
: 2026-10-04 15:43:22.369 INFO Chat Model Completion:
["The weather in Katowice is currently 20°C and sunny."]
: 2026-10-04 15:43:22.369 INFO Chat Model Prompt Content:
["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowice?", "", ""]
: 2026-10-04 15:43:20.734 INFO [Advisor] BEFORE - 4 messages, user="How is the weather in Katowice?"
: 2026-10-04 15:43:20.732 INFO Spring AI Tool: getWeather(Katowice)
: 2026-10-04 15:43:20.728 INFO [Advisor] AFTER - toolCalls=true, responseLength=0
: 2026-10-04 15:43:20.718 INFO Chat Model Completion:
[]
: 2026-10-04 15:43:20.718 INFO Chat Model Prompt Content:
["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowice?"]
: 2026-10-04 15:43:19.144 INFO [Advisor] BEFORE - 2 messages, user="How is the weather in Katowice?"
: 2026-10-04 15:43:19.129 INFO Spring AI chat [advisors]: conversationId=user-f91vbd
```

Advisor Order

The execution order of advisors in the chain is determined by the `getOrder()` method.

Lower order = earlier in the chain,
farther from the model [LLM].

Higher order = later in the chain,
closer to the model [LLM].




Example of Advisors

	Order ID Value
MessageChatMemoryAdvisor → adds history as separate messages	MIN_VALUE + 200
QuestionAnswerAdvisor → adds context from a vector store (RAG)	0
SafeGuardAdvisor → checks responses for disallowed content	0
ToolCallingAdvisor → owns the tool calling lifecycle	MIN_VALUE + 300
SimpleLoggerAdvisor → logs request and response information	0


Integer.MIN_VALUE = -2 147 483 648 ← HIGHEST_PRECEDENCE (doc name)

Default



```
3e3826b1e23115f...
-| : 2026-10-04 15:39:08.988 INFO Chat Model Prompt Content:
    ["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowice?", ""
    ""]
: 2026-10-04 15:39:08.987 INFO Chat Model Completion:
    ["The weather in Katowice is currently 20°C and sunny."]
: 2026-10-04 15:39:08.070 INFO Spring AI Tool: getWeather(Katowice)
: 2026-10-04 15:39:08.050 INFO Chat Model Prompt Content:
    ["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowice?"]
: 2026-10-04 15:39:08.049 INFO Chat Model Completion:
    []
: 2026-10-04 15:39:04.499 INFO Spring AI chat [default]: conversationId=user-bogff1
```

Memory in the Loop



```
.d8bbb3d04f0166d6a...  
: 2026-10-04 15:43:53.762 INFO Chat Model Completion:  
  ["The weather in Katowice is currently 20°C and sunny."]  
: 2026-10-04 15:43:53.762 INFO Chat Model Prompt Content:  
  ["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowice?", "How  
  is the weather in Katowice?", "", ""]  
: 2026-10-04 15:43:52.928 INFO Spring AI Tool: getWeather(Katowice)  
: 2026-10-04 15:43:52.916 INFO Chat Model Prompt Content:  
  ["You are a helpful assistant. Use available tools when needed.", "How is the weather in Katowice?"]  
: 2026-10-04 15:43:52.916 INFO Chat Model Completion:  
  []  
: 2026-10-04 15:43:51.878 INFO Spring AI chat [memory in loop]: conversationId=user-gco646
```

Three levels of hiding the loop

BY HAND

raw-agent


```
while (i++ < MAX) {  
  resp = model.call(messages);  
  if (!resp.hasToolCalls()) return resp;  
  for (call : resp.toolCalls())  
    messages.add(execute(call));  
}
```



FRAMEWORK SERVICE

LangChain4j AIServices

```
interface Assistant {  
  String chat(String msg);  
}  
AIServices.builder(Assistant.class)  
  .chatModel(model).tools(tools)  
  .build();
```



ADVISOR CHAIN

Spring AI ChatClient


```
chatClient.prompt(msg)  
  .tools(tools)  
  .call()  
  .content();
```



OpenRouter

OpenRouter is a single API in front of hundreds of models from dozens of vendors.

One key, one base URL, and the OpenAI protocol everyone already speaks.



Configuration

Spring AI configuration

```
spring:
  config:
    activate:
      on-profile: openrouter
  ai:
    model:
      chat: openai
    openai:
      base-url: https://openrouter.ai/api/v1
      api-key: ${OPENROUTER_API_KEY}
      chat:
        options:
          model: ${OPENROUTER_MODEL:openai/gpt-4o-mini}
          temperature: 0.3
      image:
        options:
          model: ${OPENROUTER_IMAGE_MODEL:google/gemini-3-pro-image}
          size: 1024x1024
      audio:
        speech:
          options:
            model: ${OPENROUTER_TTS_MODEL:x-ai/grok-voice-tts-1.0}
            voice: ${OPENROUTER_TTS_VOICE:eve}
            response-format: mp3
            speed: 1.0
        transcription:
          options:
            model: ${OPENROUTER_STT_MODEL:openai/whisper-large-v3}
```

Own properties

```
multimodal:
  media-dir: ./media
  vision-model: ${OPENROUTER_VISION_MODEL:openai/gpt-4o}
  music:
    model-pro: ${OPENROUTER_MUSIC_MODEL_PRO:google/lyria-3-pro-preview}
    model-clip: ${OPENROUTER_MUSIC_MODEL_CLIP:google/lyria-3-clip-preview}
  video:
    enabled: ${MULTIMODAL_VIDEO_ENABLED:false}
    model: ${OPENROUTER_VIDEO_MODEL:google/veo-3.1}
    resolution: ${OPENROUTER_VIDEO_RESOLUTION:720p}
```

Chat

```
public String chat(String conversationId, String question, String imageUrl) {
    String message = (imageUrl == null || imageUrl.isBlank())
        ? question
        : question + "\n\n[The user attached an image at " + imageUrl + "];

    log.info("[MM] router conversationId={} image={} question='{}'",
        conversationId, imageUrl, question);

    List<Object> toolBeans = new ArrayList<>(List.of(tools, musicTools));
    if (videoTools != null) {
        toolBeans.add(videoTools);
    }

    return router.prompt() ChatClientRequestSpec
        .advisors(MessageChatMemoryAdvisor.builder(chatMemory).build())
        .advisors(a -> a.param(ChatMemory.CONVERSATION_ID, conversationId))
        .tools(toolBeans.toArray())
        .user(message)
        .call() CallResponseSpec
        .content();
}
```

Image Generation

```
try {
    ImageResponse response = imageModel.call(new ImagePrompt(description));
    var image = response.getResult().getOutput();

    byte[] bytes;
    if (image.getB64Json() != null && !image.getB64Json().isBlank()) {
        bytes = Base64.getDecoder().decode(image.getB64Json());
    } else if (image.getUrl() != null && !image.getUrl().isBlank()) {
        bytes = download(image.getUrl());
    } else {
        collector.recordFailure();
        record(modality: "image", bytes: 0, sample);
        return "ERROR: the image model returned neither b64_json nor a url. Do NOT retry.";
    }
}
```

Speech

```
public String speak_text(  
    @ToolParam(description = "The exact text to read out loud, in the language it is written in")  
    String text) {  
  
    String refusal = refuseIfBudgetSpent(tool: "speak_text", capability: "read text out loud");  
    if (refusal != null) {  
        return refusal;  
    }  
  
    log.info("[MM] speak_text: {} chars", text == null ? 0 : text.length());  
    Timer.Sample sample = Timer.start(registry);  
    try {  
        byte[] mp3 = speechModel.call(text);  
        String url = storage.store(mp3, extension: "mp3");  
        String scriptUrl = storage.store(text.getBytes(StandardCharsets.UTF_8), extension: "txt");  
        collector.add(Artifact.of(scriptUrl, Artifact.Kind.TEXT, abbreviate(text)));  
        collector.add(Artifact.of(url, Artifact.Kind.SPEECH, abbreviate(text)));  
  
        record(modality: "speech", mp3.length, sample);  
        return "Done — the audio has been generated and is playing for the user.";  
    } catch (RuntimeException e) {  
        int failures = collector.recordFailure();  
        log.error("[MM] speak_text failed ({} / {}): {}",  
            failures, MediaCollector.MAX_FAILURES, e.getMessage());  
        record(modality: "speech", bytes: 0, sample);  
        return "ERROR: could not synthesize speech — " + e.getMessage()  
            + ". Do NOT retry; tell the user this part failed.";  
    }  
}
```

Vision

```
public String analyze(String imageUrl, String question) {
    String relative = imageUrl.replace(MediaStorage.URL_PREFIX, replacement: "");
    Path file = storage.resolve(relative);
    log.info("[MM] analyze_image {} question='{}'", relative, question);

    Timer.Sample sample = Timer.start(registry);
    try {
        MimeType mime = relative.endsWith(".png")
            ? MimeTypeUtils.IMAGE_PNG : MimeTypeUtils.IMAGE_JPEG;

        String answer = vision.prompt() ChatClientRequestSpec
            .options(OpenAiChatOptions.builder().model(model))
            .user(u -> u
                .text(question == null || question.isBlank()
                    ? "Describe this image."
                    : question)
                .media(new Media(mime, new FileSystemResource(file))))
            .call() CallResponseSpec
            .content();

        record(sample);
        return answer;
    } catch (RuntimeException e) {
        log.error("[MM] analyze_image failed: {}", e.getMessage());
        record(sample);
        throw e;
    }
}
```

Music

Here is where the trouble started....

Lyria is a chat endpoint. Getting the bytes out took five attempts.

1. `chatModel.call()` → timed lyrics, no audio “[0.0:7.3] O-Zone, we dance tonight...”
2. `+ getMedia()` → media blocks = 0
3. `+ modalities` → `voice` is required, but was not set
4. `+ voice` → 400: Audio output requires stream: true
5. `+ stream()` → the stream arrives, media blocks = 0

Every one of those errors was correct and specific. None of them said where to look.

Why it's impossible for Spring AI?

Layer	Behavior
Google	Exposes Lyria through the Interactions API, not /audio/speech
OpenRouter	Returns audio from chat only when stream: true
OpenAI SDK	Validates a speech-oriented request schema and requires voice
Spring AI	Maps audio to Media only in the non-blocking path

No layer made a mistake, so no layer could warn.

Music dropped to raw HTTP. One modality on the transport is cheaper than pretending the abstraction covers five.

Code Snippet


```
public Result generate(String model, String prompt) {
    Map<String, Object> body = new LinkedHashMap<>();
    body.put("model", model);
    body.put("messages", List.of(Map.of(k1: "role", v1: "user", k2: "content", prompt)));
    body.put("modalities", List.of("text", "audio"));
    body.put("audio", Map.of(k1: "format", v1: "mp3"));
    body.put("stream", true);

    HttpRequest request = HttpRequest.newBuilder(URI.create(baseUrl + "/chat/completions"))
        .header(name: "Authorization", value: "Bearer " + apiKey)
        .header(name: "Content-Type", value: "application/json")
        .header(name: "Accept", value: "text/event-stream")
        .timeout(Duration.ofMinutes(5))
        .POST(HttpRequest.BodyPublishers.ofString(mapper.writeValueAsString(body)))
        .build();

    try {
        HttpResponse<Stream<String>> response =
            http.send(request, HttpResponse.BodyHandlers.ofLines());



        if (response.statusCode() != 200) {
            throw new IllegalStateException("Lylia returned HTTP " + response.statusCode());
        }
        return consume(response.body());
    }
}
```

Bytes



Video

TO DO One provider does not mean one protocol.



See you at part II next week!

